

# **Study of parallelization of sorting algorithms**

**CSC 746, High Performance Computing**

# Problem Statement

- Focus is on optimizing sorting algorithms, specifically mergeSort and quickSort
- Sequential mergeSort and quickSort have  $O(n \cdot \log(n))$  time complexity
- Opportunity for parallelization due to the presence of independent, repetitive tasks
- Benchmarking parallelized quickSort and mergeSort against their serial counterparts
- Using the openMP (Open Multi-Processing) API

# Codebase

- **benchmark.cpp** -> driver program which does test runs, creates problem sizes and runs the different algorithms for each problem size and measures the runtime
- **setUpArray.cpp** -> initialises array with random numbers of problem size
- **mergeSort & quickSort.cpp** -> sequential version of the sorting algorithms
- **mergeSort\_openMP & quickSort\_openMP.cpp** -> parallelized openMP versions of the sorting algorithms

# Methodology

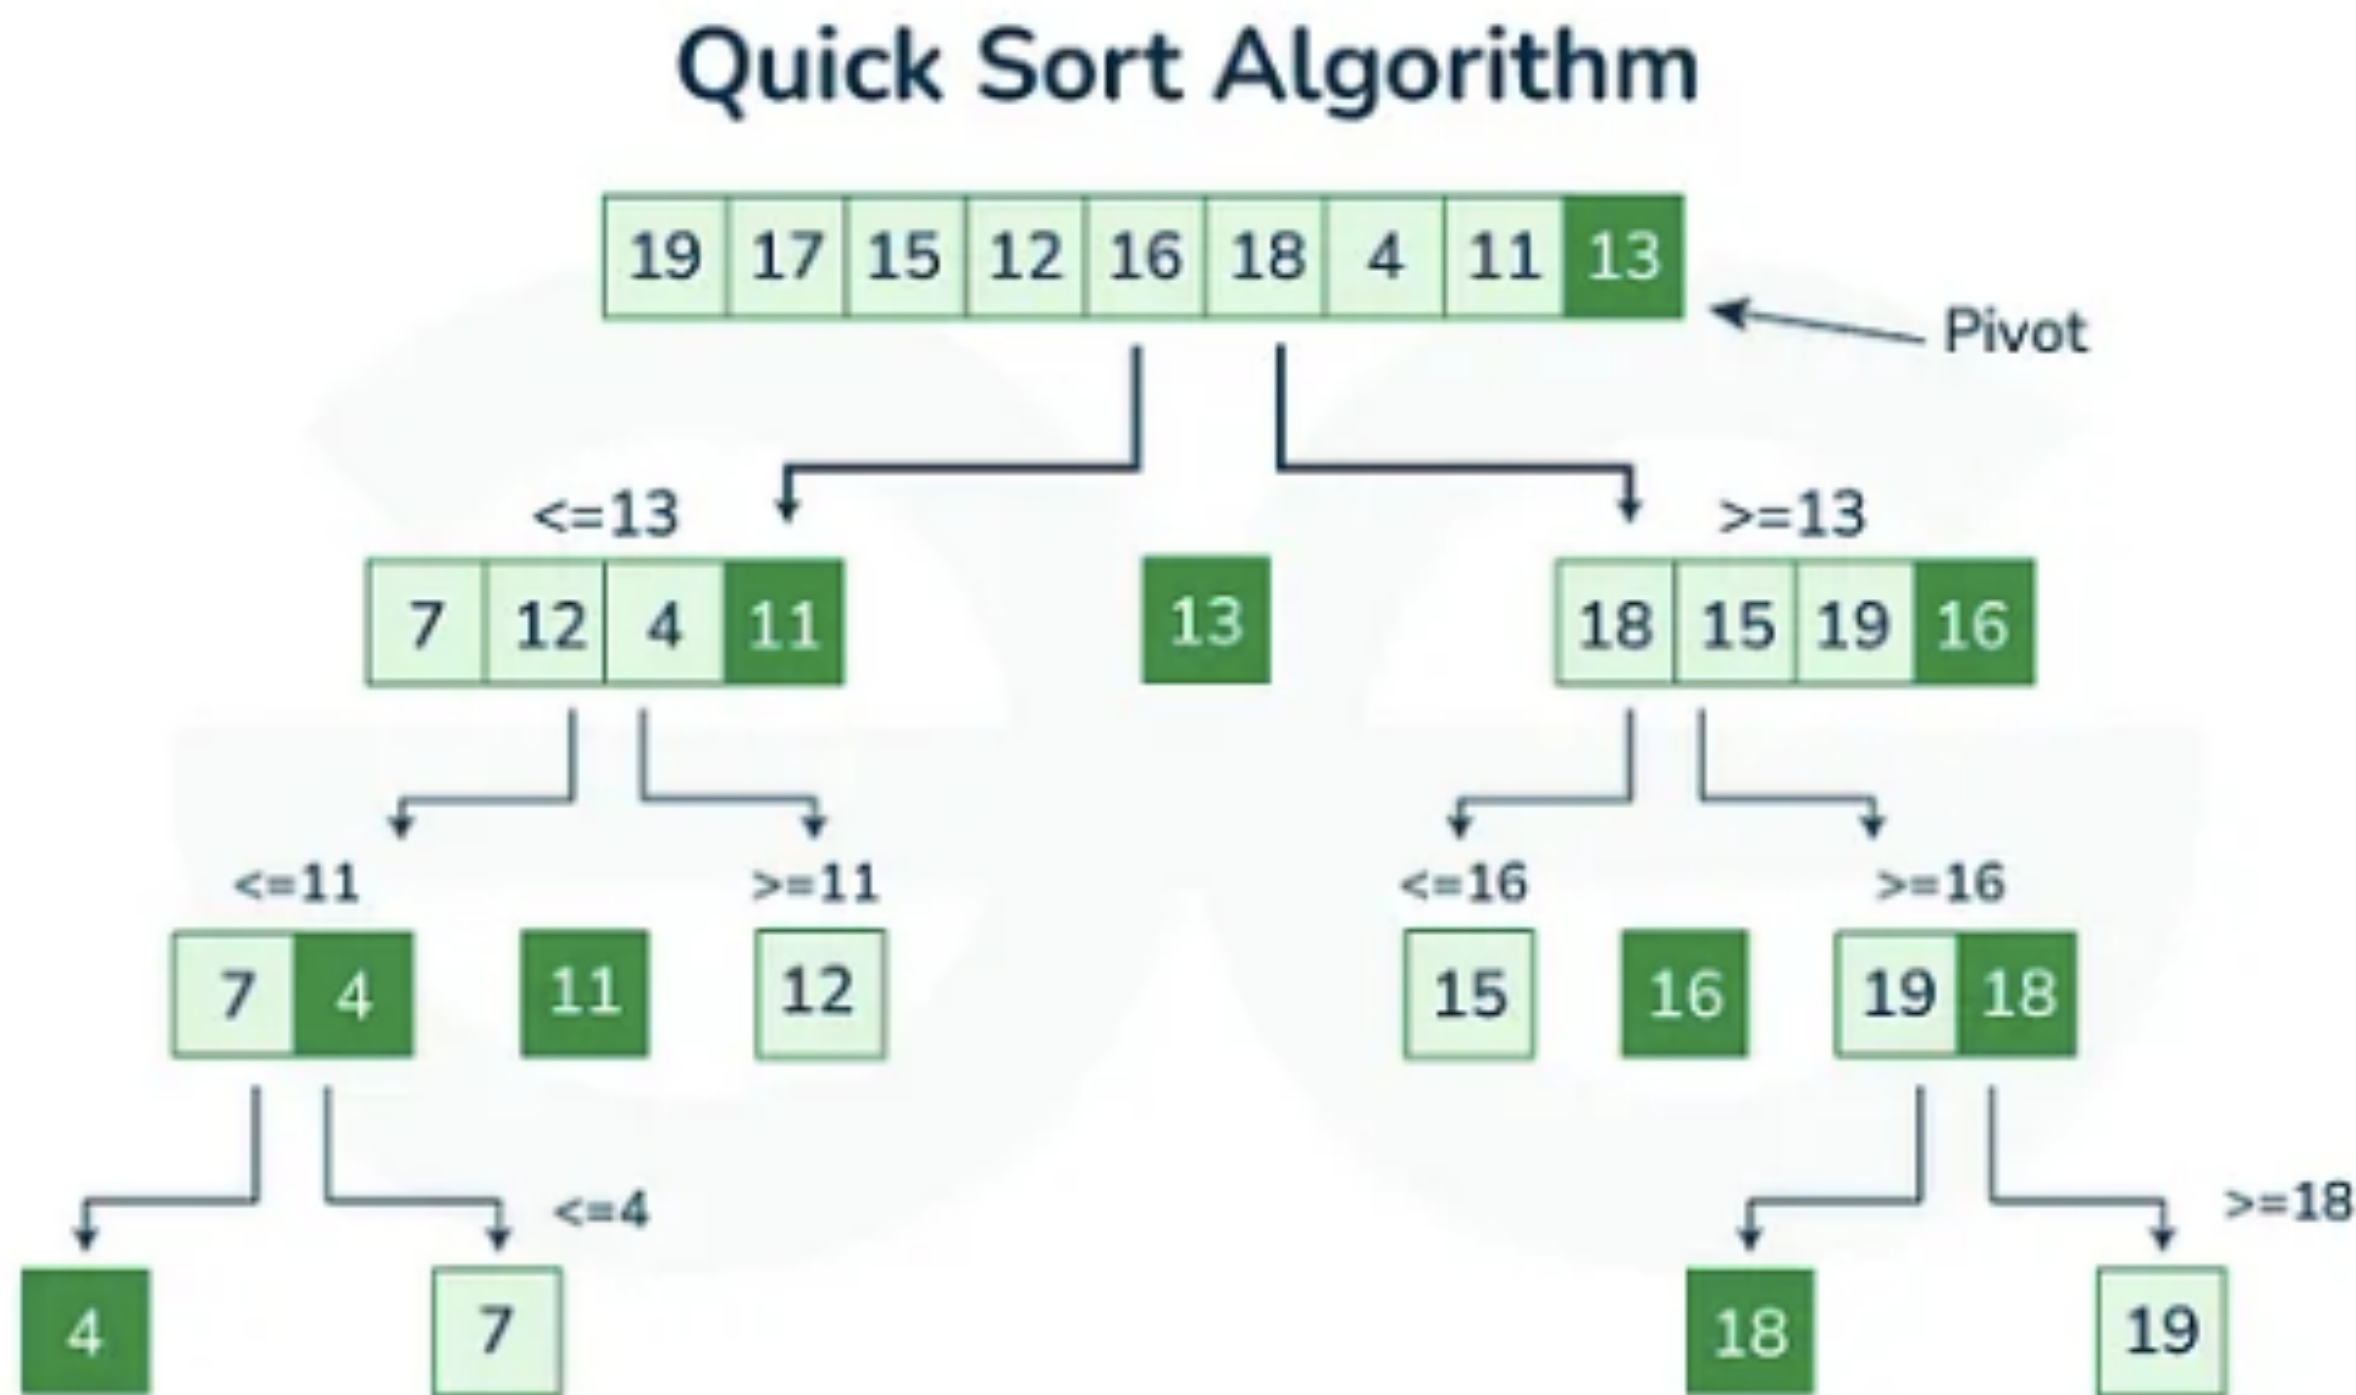
- two different performance metrics were calculated using elapsed runtime
- runtime was measured by adding instrumentation code around the sorting methods (C++ Chrono Timer)
- ➡ runtime across problem sizes and concurrency level
- ➡ speed up across concurrency levels
  - $\text{speedup}(x) = (\text{runtime\_at\_concurrency\_1}) / (\text{runtime\_at\_concurrency\_x})$
- tests were conducted over 6 different concurrency levels, ranging from 1, 4, 9, 15, 40, 60 with problem sizes ranging from 16, 32, 64, 128, 265 - million

# Computational environment and platform

- experiments were conducted on a single CPU node of the Perlmutter supercomputer (NERSC-9)
- a CPU node features 2 AMD EPYC 7763 (Milan) CPU's, each with a base clock speed of 2.45 GHz and 256 MB of L3 Cache and 512 GB of DDR4 memory
- OS is the Cray Linux Environment
- source code was compiled with C++ (GCC version 11.4.0)

# Implementation of quickSort

- select pivot and sort the array recording to this pivot in place



```
// lb = lowerBound
// ub = upperBound
partition(lb, ub, Array){
    pivot := Array[ub]
    i := lb

    for i to ub:
        if Array[i] < pivot:
            swap Array[i] and Array[lb]
            lb++

    swap pivot and Array[lb]
    return lb //new pivot index
}

QuickSort(lb, ub, Array){
    if lb < ub:
        pivotIndex = partition(lb, ub, Array)

        QuickSort(lb, pivotIndex-1, Array)
        QuickSort(pivotIndex+1, ub, Array)
}
```

# openMP task construct

- break up a big task into multiple smaller tasks and execute these tasks across the available threads.

```
#pragma omp task
{
    // this task gets written on a sticky note
    // and added to the list of tasks
}
```

```
#pragma omp single
{
    #pragma omp task
    // this task gets written on a sticky note
    // and added to the list of tasks

    #pragma omp task
    // this task gets written on a sticky note
    // and added to the list of tasks
}
```

- one master thread to write the tasks down to sticky notes
- all threads execute the sticky notes

sticky notes



Team of threads





# quickSort with openMP task construct

## parallelized version

- breaking down recursive calls into independent tasks
- those independent tasks can run in parallel
- openMP handles dependencies (sub arrays depending on partitioning)
- small array sizes are sorted sequential to avoid overhead
- shared (Array) = array is shared among all tasks
- threshold (taskLimit) is 300

```
void quickSort(int64_t lb, int64_t ub, uint64_t* Array, int64_t taskLimit){  
  
    if(lb < ub)  
    {  
        if ((ub - lb) < taskLimit ){  
            return quickSort_sequential(lb, ub, Array);  
        }else{  
            int64_t pivot_index = partition_openMP(lb, ub, Array);  
  
            #pragma omp task shared(Array)  
            quickSort(lb, pivot_index - 1, Array, taskLimit);  
            #pragma omp task shared(Array)  
            quickSort(pivot_index + 1, ub, Array, taskLimit);  
        }  
    }  
}
```

```
void quickSort_openMP(int64_t lb, int64_t ub, uint64_t* Array){  
    #pragma omp parallel  
    {  
        #pragma omp single  
        quicksort(lb, ub, Array, 300);  
        #pragma omp taskwait  
    }  
}
```

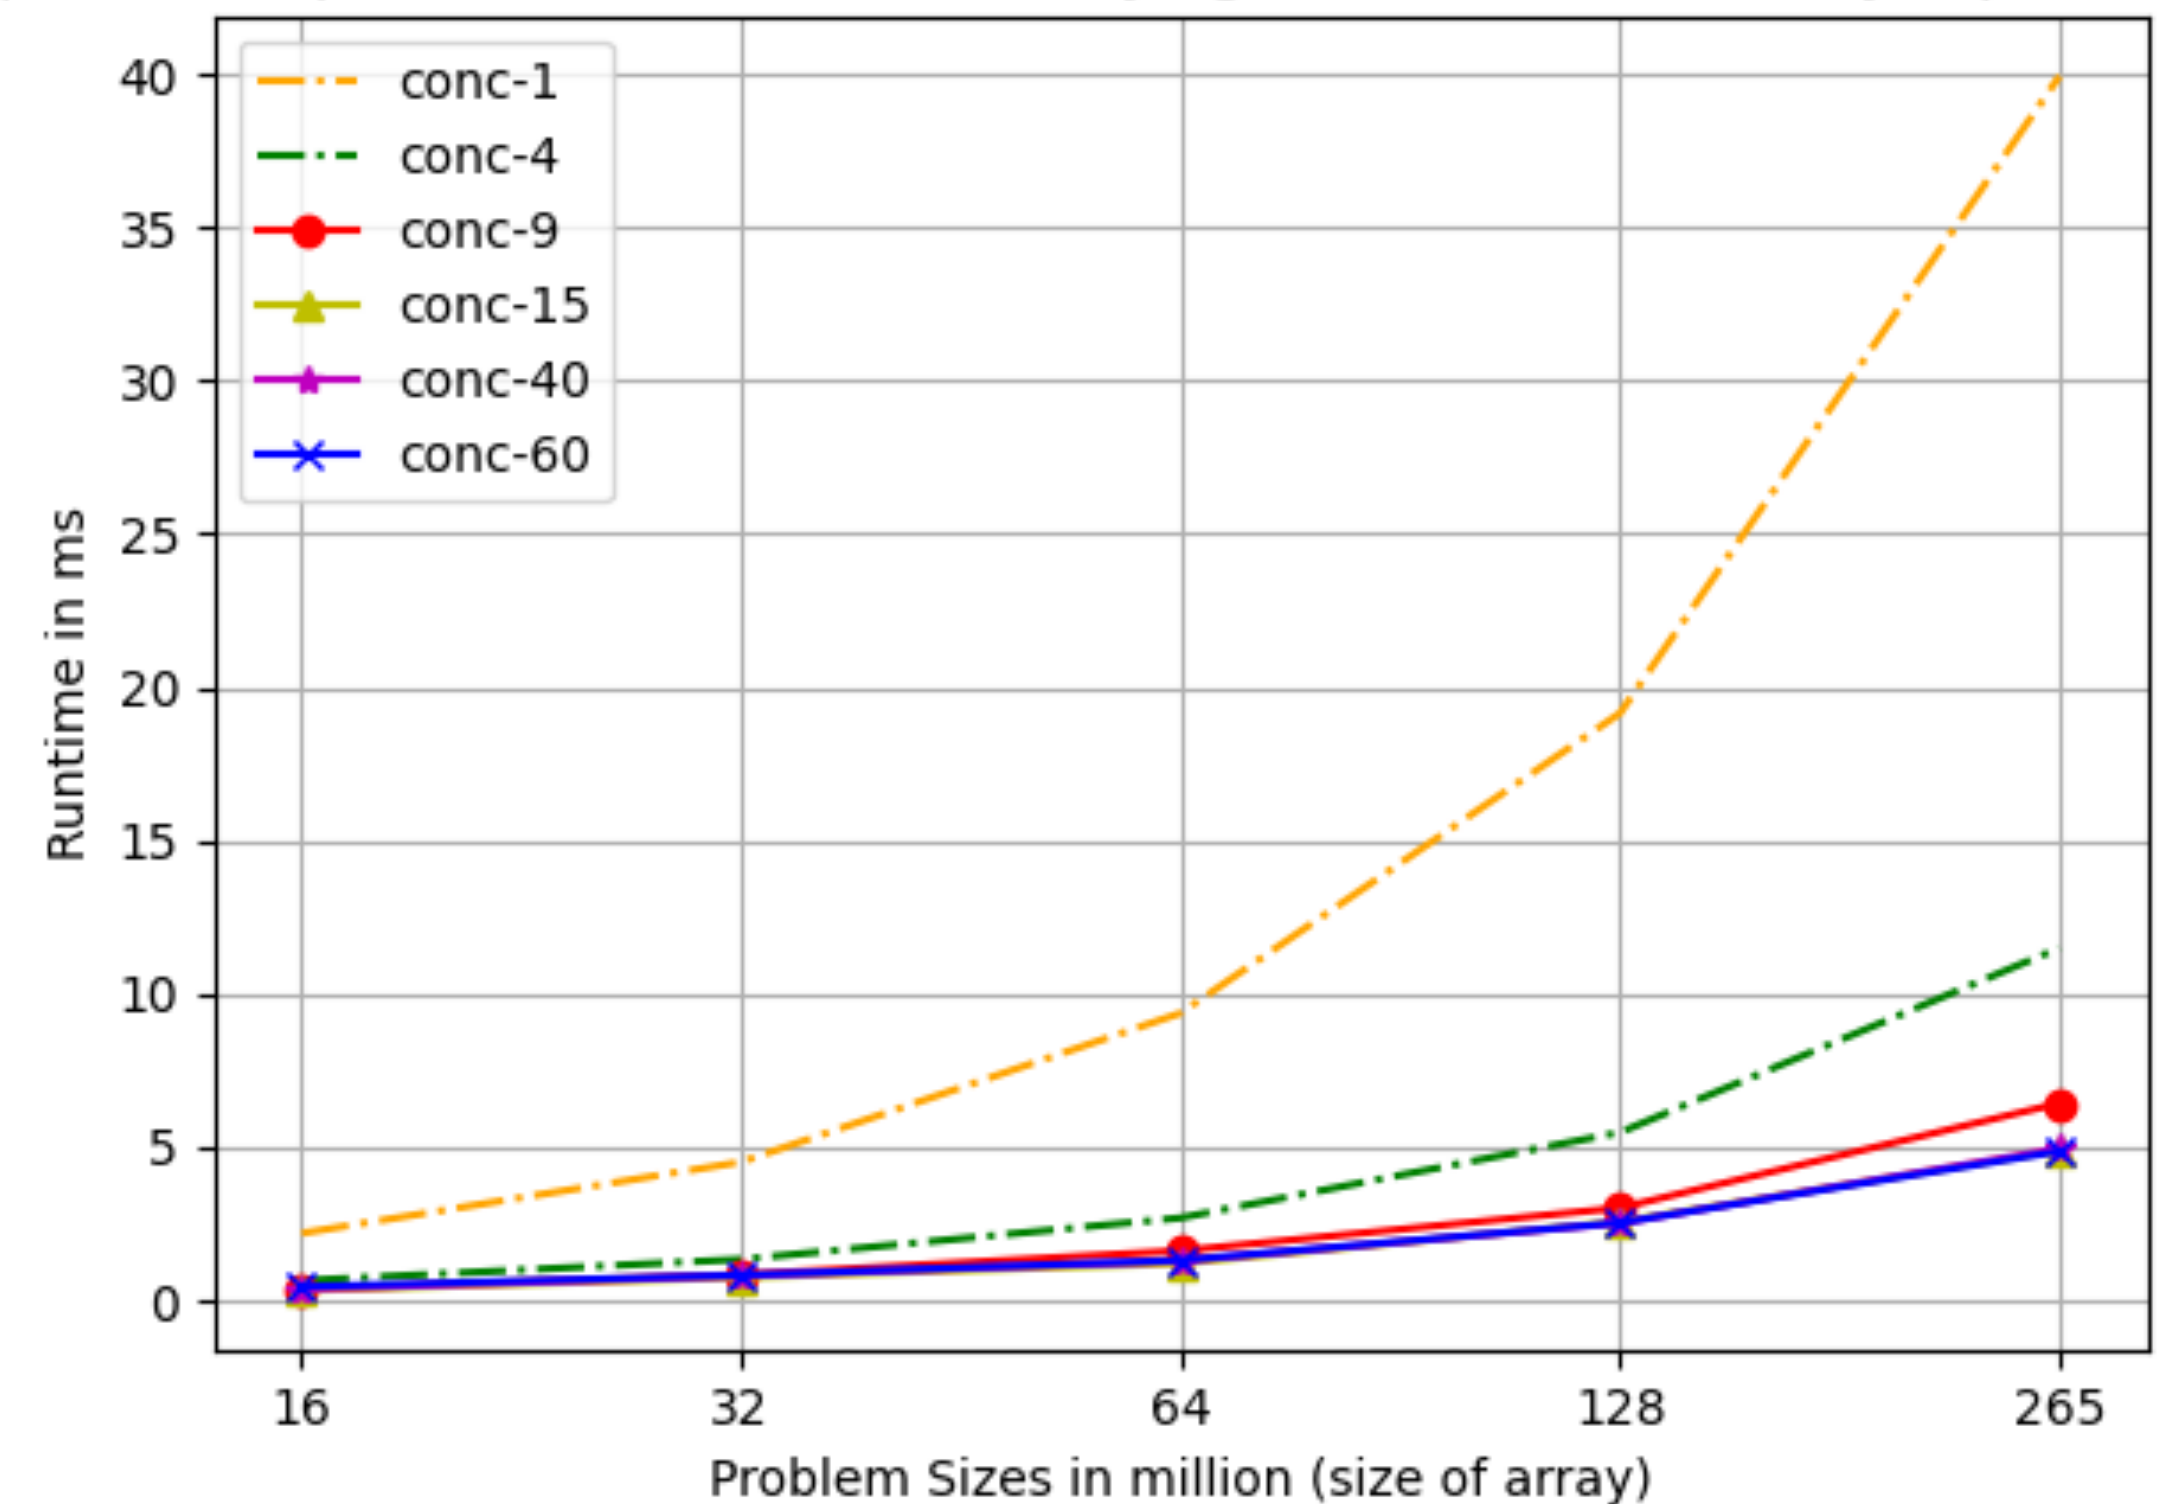


# Result: quickSort

## Runtime across problem sizes and concurrency level

- conc-1 (sequential) has highest runtimes
- all other conc-levels are very similar for smaller problem sizes
- for larger problem sizes conc-4 and conc-9 show bigger runtimes
- conc-15, conc-40 and conc-60 remain very similar

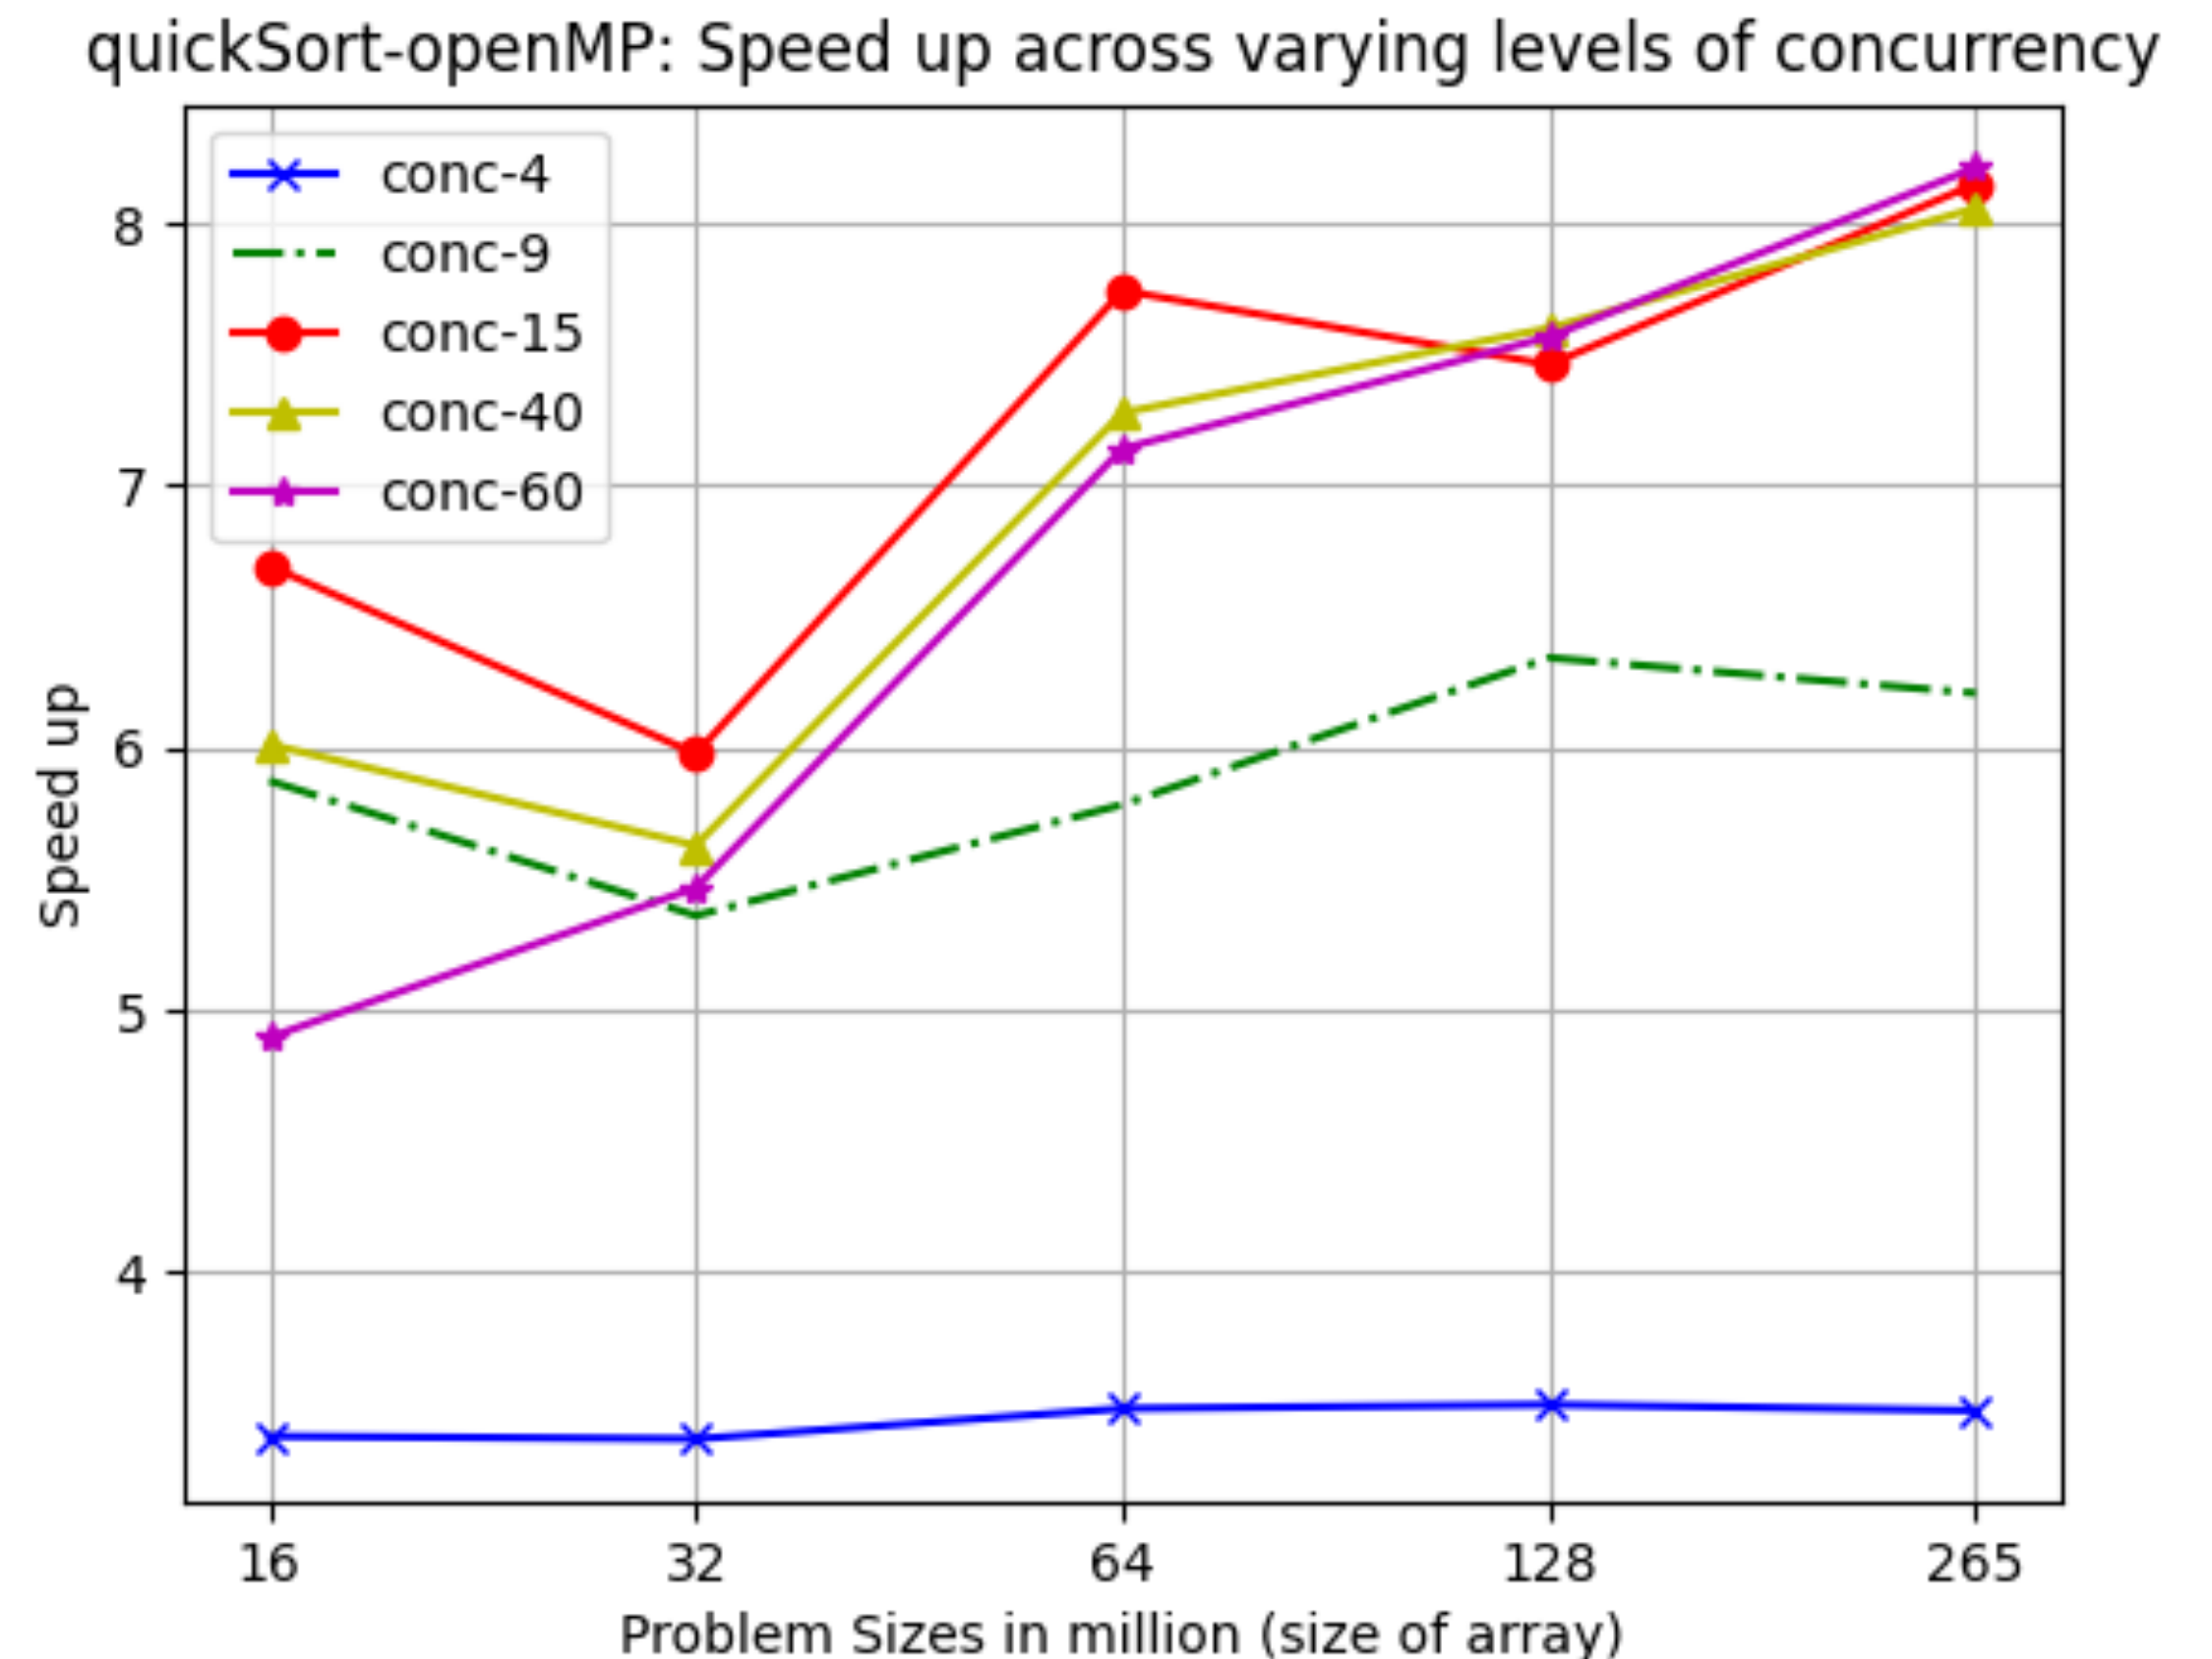
quickSort-openMP: Runtime across varying levels of concurrency & problem sizes



# Result: quickSort

## Speed up across varying levels of concurrency

- $\text{speedup}(x) = (\text{conc}-1) / (\text{conc}-x)$
- $\text{speedup}(4)$  remains steady rate of  $\sim 3.4$
- a decline in speed up for problem size 32 million
- conc-15 has the best speed up across all problem sizes
- conc-15, conc-40 and conc-60 are more than 8 times faster than the sequential version



# Findings

- significant performance improvement was achieved by parallelizing the quickSort algorithm (especially with concurrency level 15)
- openMP task construct is a good approach for parallelizing recursive sorting algorithms like quickSort
- increasing threads beyond a certain point no longer improves runtime
- minimal or no benefit when parallelizing very small arrays due to the overhead of task creation and thread assignment

**The End**